

CareerForge · Day 4 · D4-C

HireSignal — Recruiter Override Request

Student: Alia Islam | Date: June 11, 2026 | Variant: D4-C — Eligibility Exception Ambiguity

DOCUMENT 1 — AMBIGUITY RESOLUTION BRIEF

1. Founder Request

The founder wants recruiters to manually override AI scores and shortlist candidates before the pilot. Delivered as one urgent sentence — no scope, permissions, or audit requirements defined.

2. Known Facts

Known Fact	Source
HireSignal uses AI to score and rank candidates.	Scenario Package
Recruiter pilot is imminent.	Founder message
Recruiters want to shortlist candidates the AI scored low.	Recruiter message
Engineering is concerned about unlogged, unexplained overrides.	Engineering Note
No full admin permissions system exists yet.	Business Constraints
Minimal safe v1 is acceptable.	Business Constraints

3. Open Ambiguities

Ambiguity	Why It Matters	Risk If Ignored
Can recruiters shortlist ineligible candidates?	Eligibility bypass is a different risk to score preference override.	Silent compliance gap in a hiring product.
Is a reason required when overriding?	Determines if an input field and audit trail are needed.	Unexplained exceptions look arbitrary to auditors.
Who can override — all recruiters or only some?	Dictates whether a permission model is needed before code ships.	Uncontrolled access if the wrong people can override.
Should AI score stay visible after override?	Affects UI and whether AI transparency is maintained.	Hiding the score removes accountability.
Are eligibility rules hard constraints or soft guidance?	Hard = compliance risk. Soft = pilot flexibility.	Building the wrong guardrails entirely.

4. One Clarification Question

"When a recruiter shortlists a candidate who is below eligibility criteria, should the system require a mandatory written reason — or can that wait until after the pilot?"

Why this question: It answers whether eligibility rules are hard or soft, whether reason capture is needed, whether a single recruiter can act alone, and what audit coverage is required. One question, four answers.

Founder response:

| "For the pilot, just let them shortlist anyone. We can add reason fields later."

5. Bounded Assumptions

Assumption	Why Reasonable	Risk If Wrong
Eligibility rules are soft guidance for the pilot.	Founder said 'let them shortlist anyone' — a legal constraint would not be dismissed this casually.	Regulatory exposure if any criterion is a legal minimum.
Pilot scope = small, trusted recruiter group only.	Request is explicitly pre-pilot. Blast radius is contained.	Multi-tenant system would break permission assumptions entirely.
Non-blocking eligibility warning is adequate coverage.	Pilot is short and small. Warning + server log is defensible minimum.	A disputed exception may need more than a warning to defend.

6. Minimal Safe V1

V1 Behaviour	Why It Is Safe Enough
Any recruiter can manually shortlist any candidate.	Founder authorised this. Blocking it defeats the pilot.
Non-blocking eligibility warning for below-criteria candidates.	Recruiter knows they're making an exception, not a standard shortlist.
Timestamp + recruiter ID logged server-side on every shortlist.	Raw log is enough to answer 'who did this' post-pilot. No UI needed.
AI score and 'Manually Shortlisted' shown as separate labels on card.	Neither replaces the other. AI signal is not hidden or erased.

7. What Not To Build Yet

Not In V1	Why Not Now	Risk Of Building It Now
Mandatory reason / justification field	Founder rejected it for the pilot.	Over-engineers before we have pilot feedback.
Hiring manager approval workflow	Introduces a stakeholder dependency that doesn't exist yet.	Delays the pilot. Wrong scope entirely.
Audit dashboard	Explicitly excluded in business constraints.	Scope creep. Raw log covers the pilot period.
Bulk override	Not requested. AI suggestion only.	Mass exceptions with zero accountability.
Candidate notifications	Requires a comms policy decision outside this sprint.	Confusing signals to candidates during a system test.

8. Minimal Interface Expectation

Field	Detail
Action	Recruiter clicks 'Shortlist' on a candidate card.
Required fields	None. One-click confirmation.
Visible result	Card shows 'Manually Shortlisted' label. AI score remains visible alongside it.
Warning state	Non-blocking modal if candidate is below eligibility criteria.

Audit note

Timestamp + recruiter ID logged server-side. No dashboard in v1.

9. AI / Founder Assumption Challenge

Founder assumed 'add it quickly' means start building. I challenged this — the request had an eligibility exception question buried inside it that carries compliance risk. One clarification question is faster than a rework.

AI recommended a full override management module: role-based permissions, approval workflow, audit dashboard, bulk override, candidate notifications, analytics. Rejected in full. The AI optimised for completeness, not pilot safety. None of it is needed for a two-week pilot with a small recruiter cohort.

10. Founder-Ready Decision

PROCEED CONDITIONALLY

Before engineering picks this up, I want to make sure we're aligned on what we're actually building.

Recruiters get a shortlist button on each candidate card. They can't save it without selecting a reason — referred by manager, culture fit, needs portfolio review, or other. That reason, their name, and a timestamp gets logged in the backend quietly. The original AI score stays visible right next to the manually shortlisted tag. Nothing gets hidden.

Candidates hear nothing. Since shortlisting is provisional for the pilot, we're not touching notifications at all.

We're also not building bulk override, approval flows, or a dashboard — that's all post-pilot. Right now we just need to know if the core override works.

One thing I need confirmed: can engineering get the backend log ready on the current timeline? Once that's green, we're good to start.

DOCUMENT 2 — STAKEHOLDER DECISION UPDATE

Student: Alia Islam | Variant: D4-C

Decision

☐ Proceed ☒ Proceed Conditionally ☐ Block Pending Clarification

Message To Founder

Hi — before we build the override feature, I had one question I needed answered first: when a recruiter shortlists someone below eligibility criteria, does the system need to capture a reason?

You said no for now — reason fields can wait. That's what I needed.

Here's what we're building: a shortlist button, a non-blocking eligibility warning, and a server-side log. AI score stays visible. That's it.

Not building: reason fields, approval workflows, dashboards, bulk override, or candidate notifications. The AI suggested all of those — wrong scope for a pilot.

We can move in 24 hours. Recommend reviewing exception data after week one before deciding on v2.

Key Ambiguity

Whether eligibility exceptions are silently permitted or governed by any policy or workflow. Getting this wrong embeds a compliance gap into a hiring product before the pilot even starts.

Clarification Question

"When a recruiter shortlists a candidate below eligibility criteria, should the system require a mandatory written reason — or can that be defined after the pilot?"

Safe Assumptions

Assumption	Why Safe Enough	What Could Break It
Eligibility rules are soft guidance for the pilot.	Founder explicitly said 'let them shortlist anyone.'	A post-pilot discovery that one criterion is a legal minimum.
Non-blocking warning + raw log = adequate audit coverage.	Small, short, time-limited pilot. Defensible minimum.	A disputed exception that needs a documented justification.
Session data covers recruiter identity logging.	No new infrastructure needed.	Shared logins that don't map to a single individual.

Scope Boundary

Not building: mandatory reason field, manager approval workflow, audit dashboard, candidate notifications, bulk override.

Why: Business constraints exclude most of this. The rest requires dependencies that don't exist yet. All of it should follow pilot feedback — not precede it.

DOCUMENT 3 — SAFE EXECUTION DECISION

Student: Alia Islam | Variant: D4-C

Minimal Safe V1

Shortlist button on candidate card. Non-blocking eligibility warning for below-criteria candidates. Server-side log (timestamp + recruiter ID). AI score and manual shortlist shown as separate labels. Nothing else.

Risks Prevented

Risk	How This Decision Reduces It
Permission risk	Server log captures who acted and when. Actions are no longer invisible.
Fairness risk	Eligibility warning means no recruiter can shortlist an exception without seeing it. AI score stays visible — not replaced.
Audit risk	Raw server log is enough for post-pilot review. No dashboard needed for a two-week pilot.
Overbuild risk	Reason fields, approval workflows, dashboards, bulk override, notifications — all rejected. AI recommendation rejected in full.

Rejected Scope

Mandatory reason field. Hiring manager approval workflow. Audit dashboard. Candidate notification system. Bulk override. Role-based permissions. Override analytics.

Why deferred: Founder confirmed reason capture can wait. Business constraints exclude dashboards. Everything else is an AI suggestion or future-state dependency. Build them after the pilot, not before.

Final Judgment

The founder's message was one sentence under pressure. That sentence had a compliance question inside it. I surfaced it, got an answer, and defined the smallest version that is safe to ship.

The AI's module recommendation was confidently wrong for this context — it optimised for completeness, not pilot safety. We don't need a full override system. We need a shortlist button, a warning, and a log. That's enough to run the pilot safely.

Decision: Proceed Conditionally.